

Rugularized Logistic Regression Application

Table of Contents

Rugularized Logistic Regression Application On Real Data.....	1
1. Dataset.....	1
2. Pre-processing.....	3
2.1 Label encoding.....	3
2.2 Cross-Validation.....	4
2.3 Feature standardization.....	4
3. Selecting hyper-parameters.....	4
3.1 Selecting best learning rate.....	4
3.2 Selecting best regularization rate.....	7
4. Computing accuracy.....	8
4.1 Training set accuracy.....	8
4.2 Validation set accuracy.....	8
4.3 Test set accuracy.....	9
5. Conclusion.....	9

Rugularized Logistic Regression Application On Real Data

1. Dataset

The file `wine.csv` contains the information related to red wine, various factors affecting the quality.

For a given wine, we would like to predict whether it is of good or bad quality. So the task we have to perform is a classification and knowing that in this case, the label consists of only two different classes, thus it is a binary classification.

```
clear
X = readtable('wine.csv'); % Load comma separated data
```

Warning: Column headers from the file were modified to make them valid MATLAB identifiers before creating variable names for the table. The original column headers are saved in the `VariableDescriptions` property.
Set `'VariableNamingRule'` to `'preserve'` to use the original column headers as table variable names.

```
summary(X) % Summary of data
```

Variables:

fixedAcidity: 1599×1 double

Properties:

Description: fixed acidity

Values:

Min	4.6
Median	7.9
Max	15.9

volatileAcidity: 1599×1 double

Properties:

Description: volatile acidity
Values:

Min	0.12
Median	0.52
Max	1.58

citricAcid: 1599×1 double

Properties:
Description: citric acid
Values:

Min	0
Median	0.26
Max	1

residualSugar: 1599×1 double

Properties:
Description: residual sugar
Values:

Min	0.9
Median	2.2
Max	15.5

chlorides: 1599×1 double

Properties:
Description: chlorides
Values:

Min	0.012
Median	0.079
Max	0.611

freeSulfurDioxide: 1599×1 double

Properties:
Description: free sulfur dioxide
Values:

Min	1
Median	14
Max	72

totalSulfurDioxide: 1599×1 double

Properties:
Description: total sulfur dioxide
Values:

Min	6
Median	38
Max	289

density: 1599×1 double

Properties:
Description: density
Values:

Min	0.99007
-----	---------

```
Median    0.99675
Max       1.0037
```

pH: 1599×1 double

```
Properties:
Description: pH
Values:
```

```
Min       2.74
Median    3.31
Max       4.01
```

sulphates: 1599×1 double

```
Properties:
Description: sulphates
Values:
```

```
Min       0.33
Median    0.62
Max       2
```

alcohol: 1599×1 double

```
Properties:
Description: alcohol
Values:
```

```
Min       8.4
Median    10.2
Max      14.9
```

quality: 1599×1 cell array of character vectors

```
Properties:
Description: quality
```

```
X(1:5,:) % 5 first examples of data set
```

```
ans = 5×12 table
```

...

	fixedAcidity	volatileAcidity	citricAcid	residualSugar	chlorides
1	7.4000	0.7000	0	1.9000	0.0760
2	7.8000	0.8800	0	2.6000	0.0980
3	7.8000	0.7600	0.0400	2.3000	0.0920
4	11.2000	0.2800	0.5600	1.9000	0.0750
5	7.4000	0.7000	0	1.9000	0.0760

2. Pre-processing

2.1 Label encoding

We start with encoding the labels of our classification, 'bad' quality is encoded as 0 and 'good' quality as 1.

```
[~, n] = size(X);
X.quality = categorical(X.quality);
```

```
% Seperate labels
y = (X{:, "quality"} == 'good');
X = X(:, 1:n-1);
n = n-1; % Number of features
```

2.2 Cross-Validation

The data is splitted to 3 parts as follows : 60% training set, 20% validation set and 20% test set.

```
% Training, 60% of dataset
TrainX = X{1:960, :};
Trainy = y(1:960, :);
Trainm = size(TrainX, 1); % Number of examples in training set

% Cross-Validation, 20% of dataset
ValidationX = X{961:1280, :};
Validationy = y(961:1280, :);
Validm = size(ValidationX, 1); % Number of examples in validation set

% Test, 20% of dataset
TestX = X{1281:end, :};
Testy = y(1281:end, :);
Testm = size(TestX, 1); % Number of examples in test set
```

2.3 Feature standardization

By looking at the values of features, we note that there are much differences between orders of magnitudes .
Feature standardization can make gradient descent converge much more quickly.

```
[TrainX, mu, sigma] = featureStandardize(TrainX); % Scale training set features and set them to zero mean
% Add intercept term to TrainX
TrainX = [ones(Trainm, 1) TrainX];
```

It is important to store the values used for standardization of training test - the mean value and the standard deviation - in order to be able to standardize validation and test set.

```
ValidationX = (ValidationX - mu)./sigma; % Scale validation set features and set them to zero mean
% Add intercept term to TrainX
ValidationX = [ones(Validm, 1) ValidationX];

TestX = (TestX - mu)./sigma; % Scale validation set features and set them to zero mean
% Add intercept term to TestX
TestX = [ones(Testm, 1) TestX];
```

3. Selecting hyper-parameters

3.1 Selecting best learning rate

In order to find the best learning rate, we train a few models with different learning rate in parallel. Then only we select the one that works better at the end on our validation set.

```
% Initialize theta
```

```


theta = zeros(n+1, 1);
iteration_max = 50;
lambda = 1;

% Choose some values for alpha
alpha_1 = 2;
alpha_2 = 2.5;
alpha_3 = 3;
alpha_4 = 3.5;
alpha_5 = 4;

% Train models with different learning rate in parallel
[theta_1, J1_history] = gradientDescent(TrainX, Trainy, theta, alpha_1, iteration_max,
[theta_2, J2_history] = gradientDescent(TrainX, Trainy, theta, alpha_2, iteration_max,
[theta_3, J3_history] = gradientDescent(TrainX, Trainy, theta, alpha_3, iteration_max,
[theta_4, J4_history] = gradientDescent(TrainX, Trainy, theta, alpha_4, iteration_max,
[theta_5, J5_history] = gradientDescent(TrainX, Trainy, theta, alpha_5, iteration_max,


% Plot the convergence graph
figure
plot(1:iteration_max, J1_history, '-y', 'LineWidth', 2);
hold on
plot(1:iteration_max, J2_history, '-r', 'LineWidth', 2);
plot(1:iteration_max, J3_history, '-k', 'LineWidth', 2);
plot(1:iteration_max, J4_history, '-g', 'LineWidth', 2);
plot(1:iteration_max, J5_history, '-b', 'LineWidth', 2);
legend('alpha= 2', 'alpha= 2.5', 'alpha= 3', 'alpha= 3.5', 'alpha= 4')
hold off
xlabel('Number of iterations');
ylabel('Training set cost J');

```



```
figure
plot(alpha_1, costFunction(theta_1, ValidationX, Validationy, lambda), '*r', 'LineWidth
hold on
plot(alpha_2, costFunction(theta_2, ValidationX, Validationy, lambda), '*r', 'LineWidth
plot(alpha_3, costFunction(theta_3, ValidationX, Validationy, lambda), '*r', 'LineWidth
plot(alpha_4, costFunction(theta_4, ValidationX, Validationy, lambda), '*r', 'LineWidth
plot(alpha_5, costFunction(theta_5, ValidationX, Validationy, lambda), '*r', 'LineWidth

hold off
xlabel('Alpha');
ylabel('Validation set cost J');
```



It can be noticed that the best learning rate is approximately 2.5 and with maximum iteration of 30, the curve reaches the minimum possible cost.

```
best_alpha = alpha_2;
iteration_max = 30;
```

3.2 Selecting best regularization rate

Same as finding the best learning rate, we train a few models with different regularization rate in parallel and select the one that works better on our validation set.

```
% Choose some values for lambda
```


```
lambda_1 = 0;
lambda_2 = 1;
lambda_3 = 2;
lambda_4 = 4;
lambda_5 = 8;
```

```
% Train models with different regularization rate in parallel
```

```
[theta_1, ~] = gradientDescent(TrainX, Trainy, theta, best_alpha, iteration_max, lambda_1);
[theta_2, ~] = gradientDescent(TrainX, Trainy, theta, best_alpha, iteration_max, lambda_2);
[theta_3, ~] = gradientDescent(TrainX, Trainy, theta, best_alpha, iteration_max, lambda_3);
[theta_4, ~] = gradientDescent(TrainX, Trainy, theta, best_alpha, iteration_max, lambda_4);
[theta_5, ~] = gradientDescent(TrainX, Trainy, theta, best_alpha, iteration_max, lambda_5);
```

```
% Plot
```

```
figure
plot(lambda_1, costFunction(theta_1, ValidationX, Validationy, lambda_1), '*r', 'LineWidth', 2);
hold on
plot(lambda_2, costFunction(theta_2, ValidationX, Validationy, lambda_2), '*r', 'LineWidth', 2);
plot(lambda_3, costFunction(theta_3, ValidationX, Validationy, lambda_3), '*r', 'LineWidth', 2);
plot(lambda_4, costFunction(theta_4, ValidationX, Validationy, lambda_4), '*r', 'LineWidth', 2);
plot(lambda_5, costFunction(theta_5, ValidationX, Validationy, lambda_5), '*r', 'LineWidth', 2);
hold off
xlabel('Lambda');
ylabel('Validation cost J');
```



It can be noticed that the lowest cost on the validation set is when regularization rate is equal to 0. In other words, fitting the parameters without regularization.

```
best_lambda = lambda_1;
theta = theta_1;
```

4. Computing accuracy

4.1 Training set accuracy

```
p = predict(theta, TrainX);
fprintf('Training Accuracy: %f\n', mean(double(p == Trainy)) * 100);
```

Training Accuracy: 74.062500

4.2 Validation set accuracy

```
p = predict(theta, ValidationX);
```



```
fprintf('Validation Accuracy: %f\n', mean(double(p == Validationy)) * 100);
```

Validation Accuracy: 78.437500

4.3 Test set accuracy

```
p = predict(theta, TestX);  
fprintf('Test Accuracy: %f\n', mean(double(p == Testy)) * 100);
```

Test Accuracy: 74.294671

5. Conclusion

Knowing that we don't notice a considerable difference between the training set accuracy and test set accuracy, it can be deduced that there is no overfitting in our model.

However, the model suffers from underfitting (high-bias + low variance). This problem would be solved by adding polynomial features or using a more complex model such as Decision Tree.